

2023/08/01

DSAIL summer-internship

STUDENT

유혜원

Hyewon Ryu

Inductive Representation Learning on Large Graphs

Table of
Contents

Page

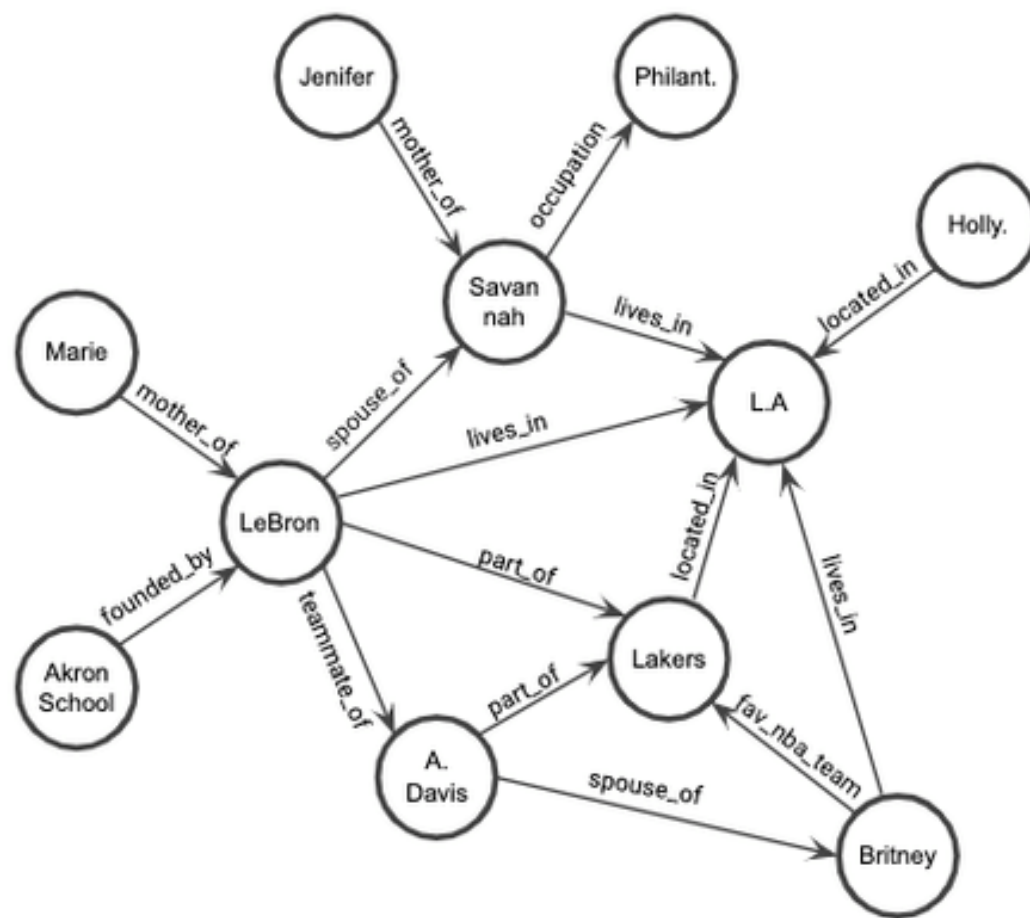
I	Research Background & Motivation	3
II	Related Works	5
III	Model Description	6
IV	Experiment Results	11
V	Implementation	13

I Research Background & Motivation

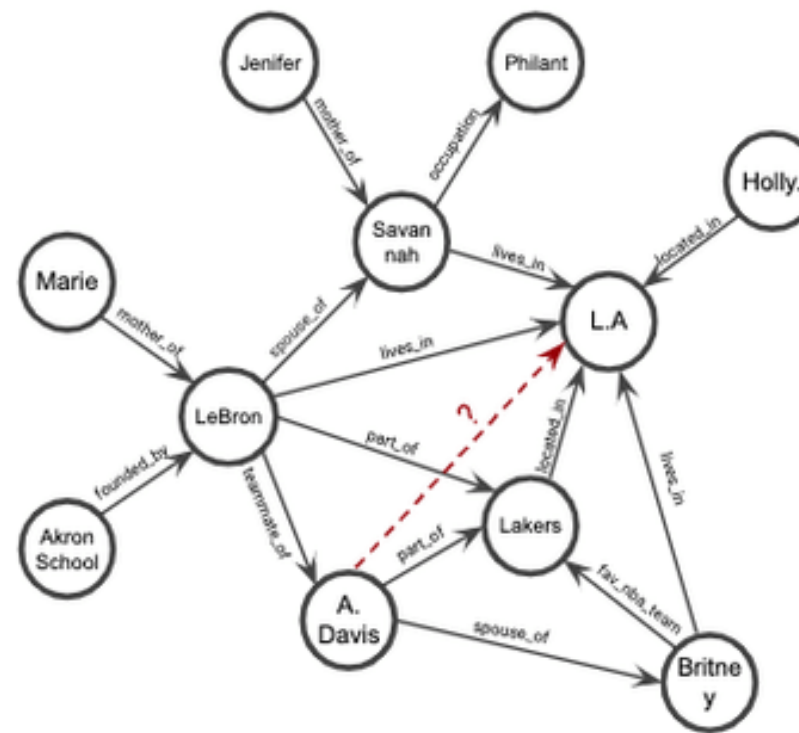
- Low-dimensional vector embedding of nodes in large graphs have proved extremely useful as feature inputs for a wide variety of prediction and graph analysis tasks
 - Node Classification, Clustering, Link Prediction
- Previous works have focused on embedding nodes from a single fixed graph
 - However, many real-world applications require quick embedding for unseen nodes, or entirely new (sub)graphs
 - We need a system which can be used on evolving graphs
 - ex) posts on Reddit, users and videos on Youtube

I Research Background & Motivation


Transductive vs. Inductive



a. Training graph



b. Transductive inference



c. Inductive inference

Most existing approaches are inherently transductive.

- Matrix-factorization-based
 - do not generalize to unseen data
- Graph Convolution Networks (GCNs) have only been applied in the transductive setting

Inductive node embedding is difficult

- Unseen nodes requires “aligning” newly observed subgraphs
- Need to learn graph structure both in local and global view

In this work, extend GCNs to the task of inductive unsupervised learning

II Related Works

1) Factorization-based embedding approaches

- Embedding using random walk statistics and matrix-factorization-based learning
 - inherently transductive
- Planetoid-I
 - inductive
 - do not use any graph structural information & feature information

2) Supervised learning over graphs

- Neural Network approaches to supervised learning over graph structures
 - main focus is to classify entire graphs, not to generate representations for individual nodes

3) Graph convolution networks

- do not scale to large graphs or are designed for whole-graph classification
- only applied to transductive setting

III Model Description

Process of Embedding Generation

Assume that the model has already been trained and parameters are fixed

→ Aggregate function parameters & weight matrix W

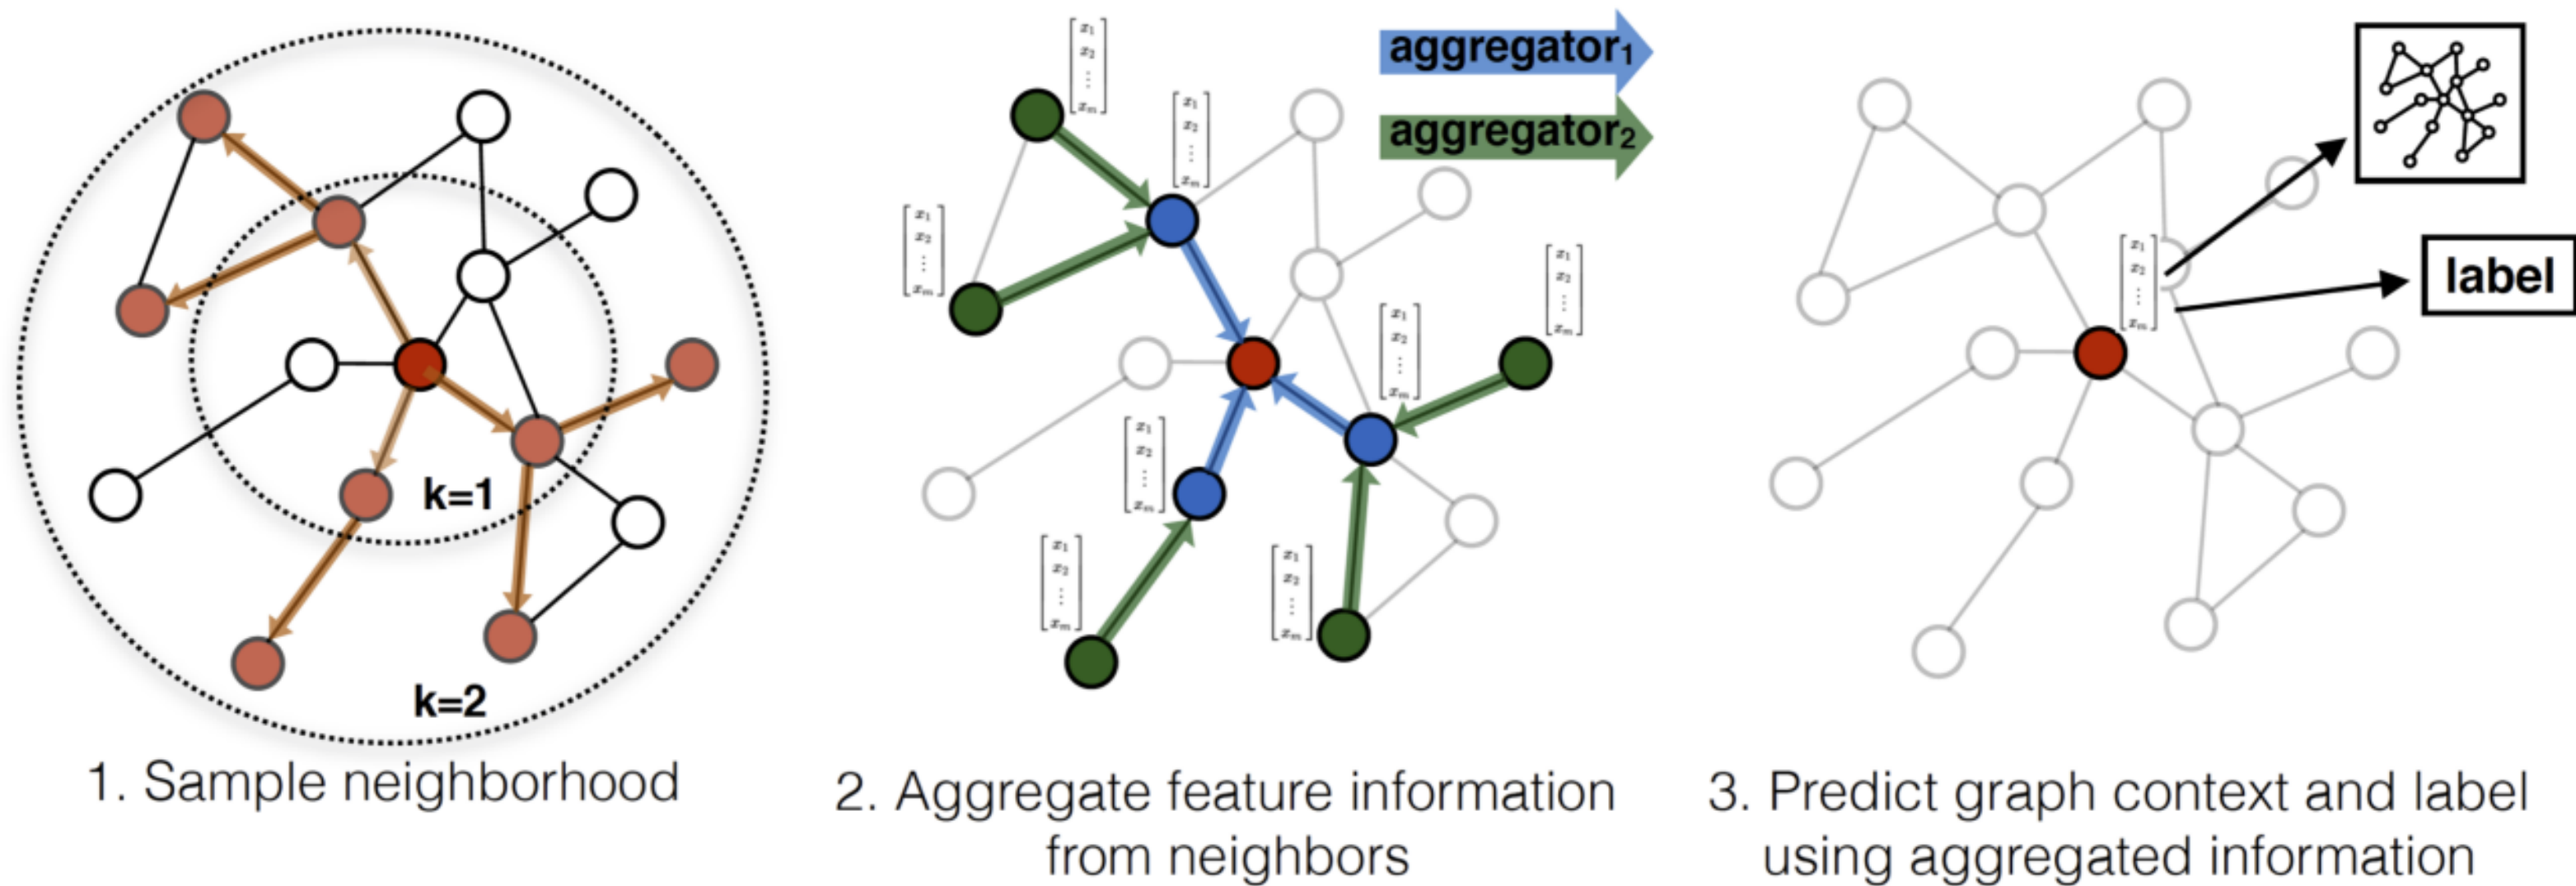
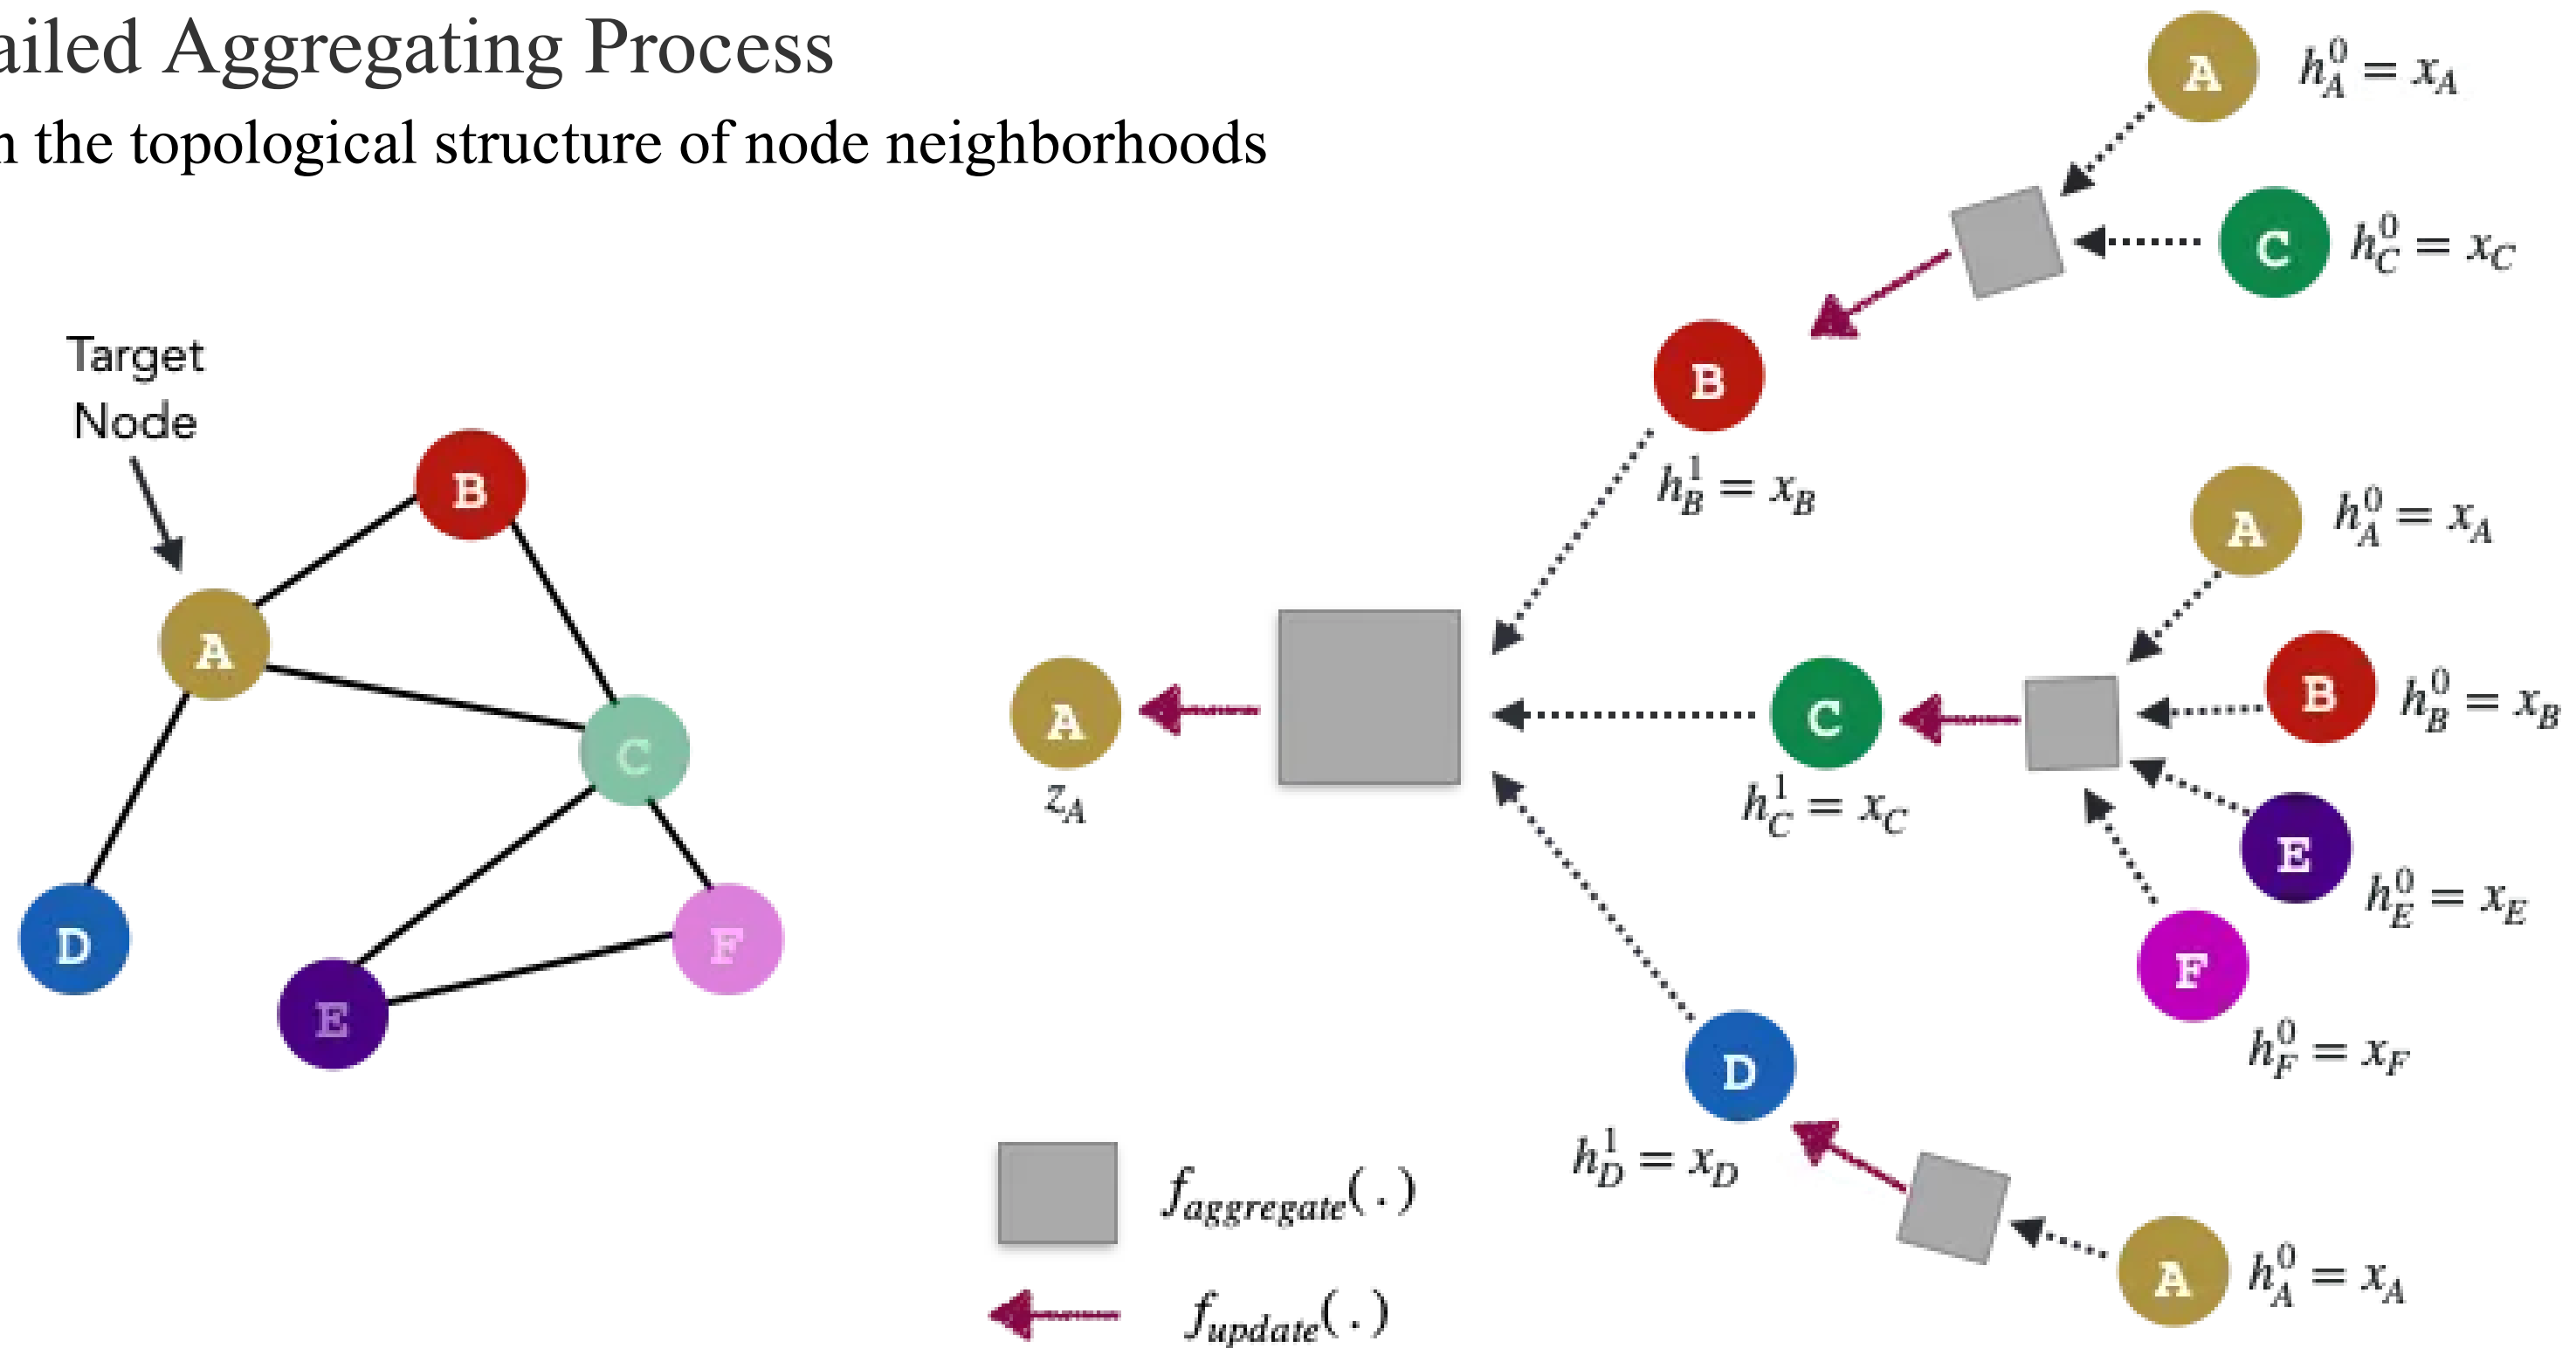


Figure 1: Visual illustration of the GraphSAGE sample and aggregate approach.

III Model Description

Detailed Aggregating Process

Learn the topological structure of node neighborhoods



III Model Description

Pseudo Code of Embedding Generation

Algorithm 1: GraphSAGE embedding generation (i.e., forward propagation) algorithm

Input : Graph $\mathcal{G}(\mathcal{V}, \mathcal{E})$; input features $\{\mathbf{x}_v, \forall v \in \mathcal{V}\}$; depth K ; weight matrices $\mathbf{W}^k, \forall k \in \{1, \dots, K\}$; non-linearity σ ; differentiable aggregator functions $\text{AGGREGATE}_k, \forall k \in \{1, \dots, K\}$; neighborhood function $\mathcal{N} : v \rightarrow 2^{\mathcal{V}}$

Output: Vector representations \mathbf{z}_v for all $v \in \mathcal{V}$

```
1  $\mathbf{h}_v^0 \leftarrow \mathbf{x}_v, \forall v \in \mathcal{V}$ ; //input node feature
2 for  $k = 1 \dots K$  do
3   for  $v \in \mathcal{V}$  do
4      $\mathbf{h}_{\mathcal{N}(v)}^k \leftarrow \text{AGGREGATE}_k(\{\mathbf{h}_u^{k-1}, \forall u \in \mathcal{N}(v)\})$ ;
5      $\mathbf{h}_v^k \leftarrow \sigma(\mathbf{W}^k \cdot \text{CONCAT}(\mathbf{h}_v^{k-1}, \mathbf{h}_{\mathcal{N}(v)}^k))$  //agreggating
6   end
7    $\mathbf{h}_v^k \leftarrow \mathbf{h}_v^k / \|\mathbf{h}_v^k\|_2, \forall v \in \mathcal{V}$ 
8 end
9  $\mathbf{z}_v \leftarrow \mathbf{h}_v^K, \forall v \in \mathcal{V}$  //final representations output at depth K
```

III Model Description

Learning the Parameters of GraphSAGE

Aggregate function parameters & weight matrix W

Definition of neighborhood

$\mathcal{N}(v)$: fixed-size, uniform draw from the set $\{u \in \mathcal{V} : (u, v) \in \mathcal{E}\}$

Loss function

- Fully unsupervised setting

$$J_{\mathcal{G}}(\mathbf{z}_u) = -\log(\sigma(\mathbf{z}_u^{\top} \mathbf{z}_v)) - Q \cdot \mathbb{E}_{v_n \sim P_n(v)} \log(\sigma(-\mathbf{z}_u^{\top} \mathbf{z}_{v_n}))$$

v : node that co-occurs near u on fixed length random walk

σ : sigmoid function

P_n : negative sampling distribution

Q : number of negative samples

- Task-specific objective
 - cross-entropy loss etc

III Model Description

Aggregator Architectures

→ Aggregator functions must operate over an unordered set of vectors

Mean Aggregator: inductive variant of the GCN

$$\mathbf{h}_v^k \leftarrow \sigma(\mathbf{W} \cdot \text{MEAN}(\{\mathbf{h}_v^{k-1}\} \cup \{\mathbf{h}_u^{k-1}, \forall u \in \mathcal{N}(v)\}))$$

→ concatenation : simple form of a “skip connection”
between the different search depth


LSTM Aggregator

- have larger expressive capability
- not symmetric
 - simply applying the LSTMs to a random permutation of the node's neighbors

Pooling Aggregator

$$\text{AGGREGATE}_k^{\text{pool}} = \max(\{\sigma(\mathbf{W}_{\text{pool}} \mathbf{h}_{u_i}^k + \mathbf{b}), \forall u_i \in \mathcal{N}(v)\})$$

Skip Connection
: prevent gradient vanishing (exploding)



IV Experiment Results

Table 1: Prediction results for the three datasets (micro-averaged F1 scores). Results for unsupervised and fully supervised GraphSAGE are shown. Analogous trends hold for macro-averaged scores.

Name	Citation		Reddit		PPI	
	Unsup. F1	Sup. F1	Unsup. F1	Sup. F1	Unsup. F1	Sup. F1
Random	0.206	0.206	0.043	0.042	0.396	0.396
Raw features	0.575	0.575	0.585	0.585	0.422	0.422
DeepWalk	0.565	0.565	0.324	0.324	—	—
DeepWalk + features	0.701	0.701	0.691	0.691	—	—
GraphSAGE-GCN	0.742	0.772	0.908	0.930	0.465	0.500
GraphSAGE-mean	0.778	0.820	0.897	0.950	0.486	0.598
GraphSAGE-LSTM	0.788	0.832	0.907	0.954	0.482	0.612
GraphSAGE-pool	0.798	0.839	0.892	0.948	0.502	0.600
% gain over feat.	39%	46%	55%	63%	19%	45%

IV Experiment Results

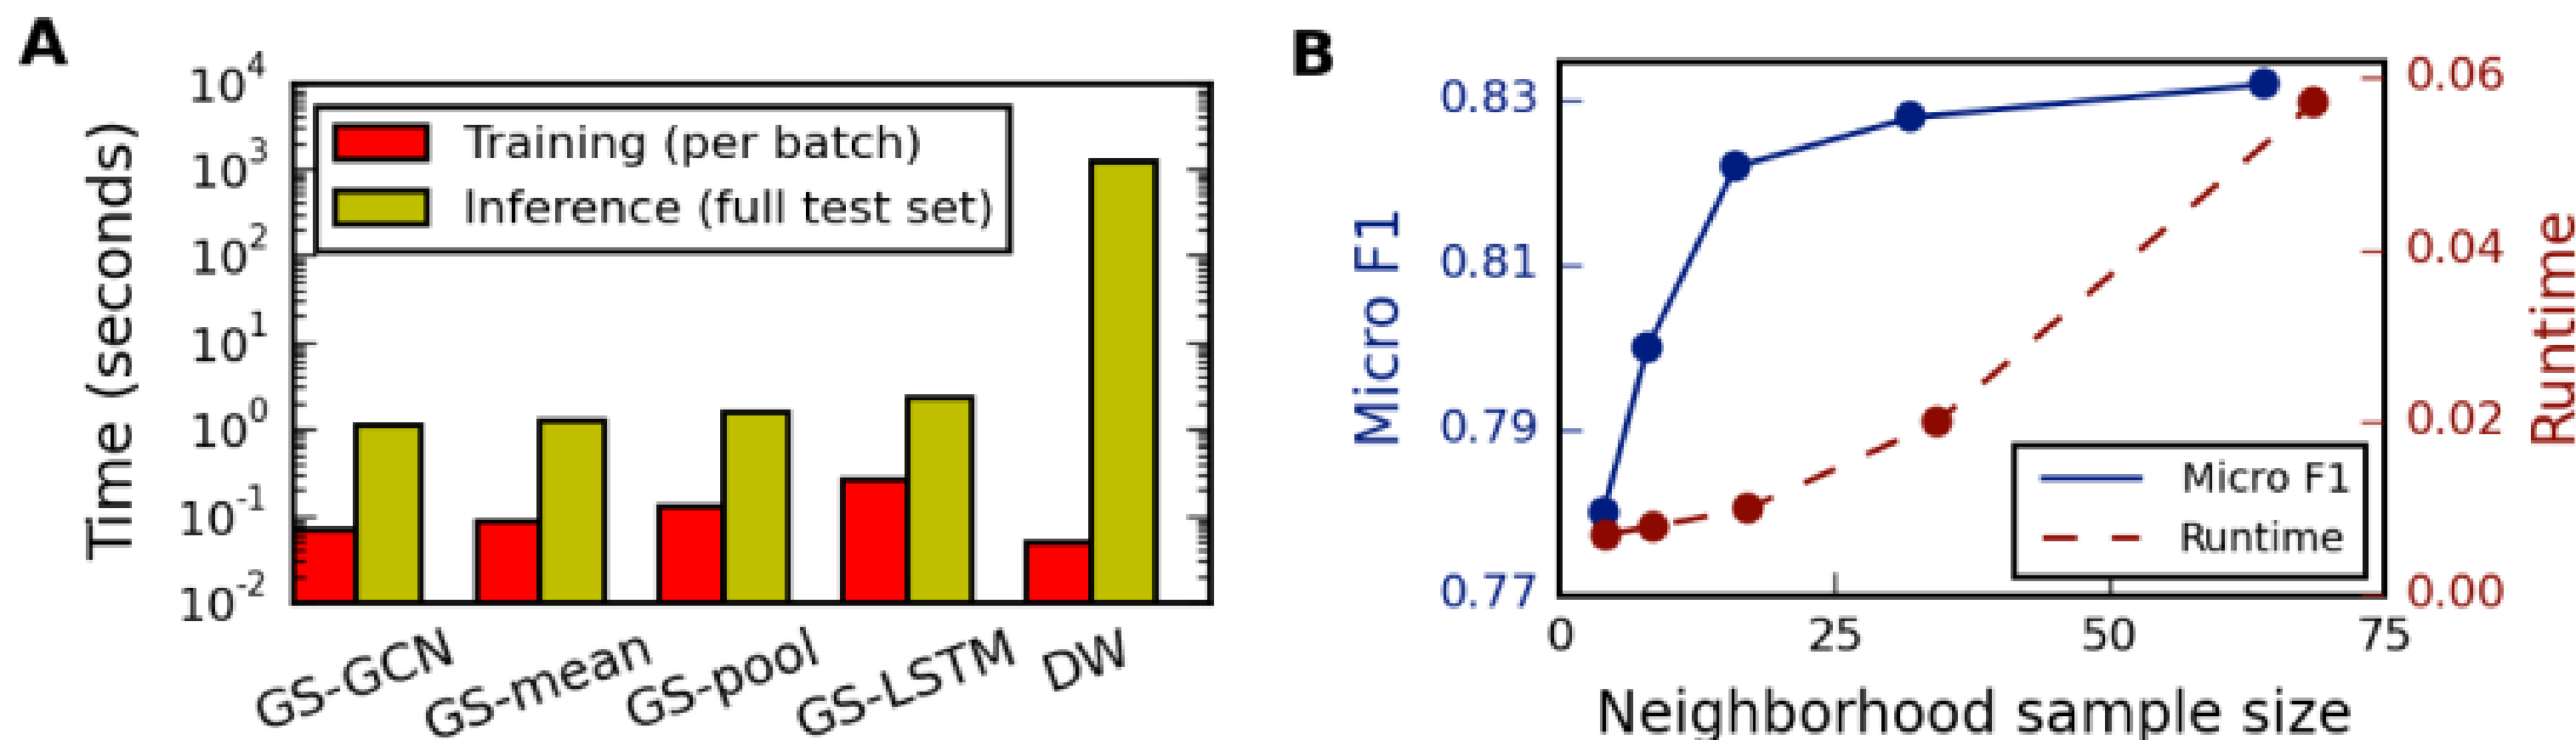


Figure 2: **A:** Timing experiments on Reddit data, with training batches of size 512 and inference on the full test set (79,534 nodes). **B:** Model performance with respect to the size of the sampled neighborhood, where the “neighborhood sample size” refers to the number of neighbors sampled at each depth for $K = 2$ with $S_1 = S_2$ (on the citation data using GraphSAGE-mean).

V Implementation

Dataset: Cora [\[not in paper\]](#)

Node: Papers in machine learning domain

Edge: Citation between papers

→ [node classification into specific domain \(7 classes\)](#)

number of nodes: 2708

number of edges: 10556

train set: 140 nodes

validation set: 500 nodes


test set: 1000 nodes

```
dataset = Planetoid(root='tmp/Cora', name='Cora')
```

V Implementation

Batch creation & Neighborhood sampling

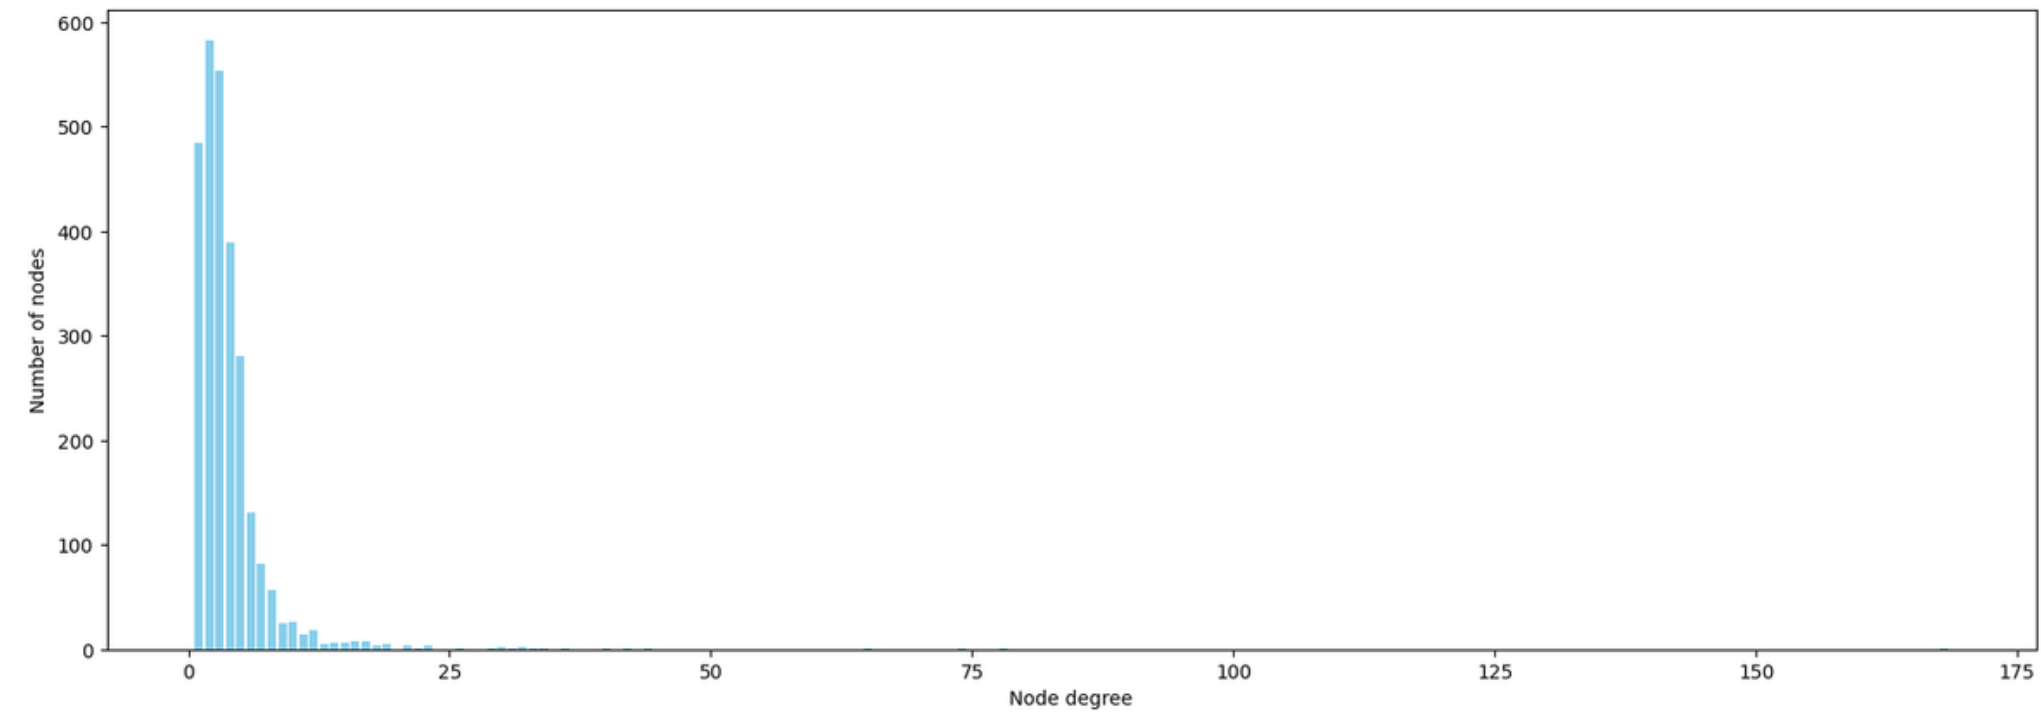
```
1 # Create batches with neighbor sampling
2 train_loader = NeighborLoader(
3     data,
4     num_neighbors = [25, 10], # same as paper
5     batch_size = 40,
6     input_nodes = data.train_mask
7 )
```



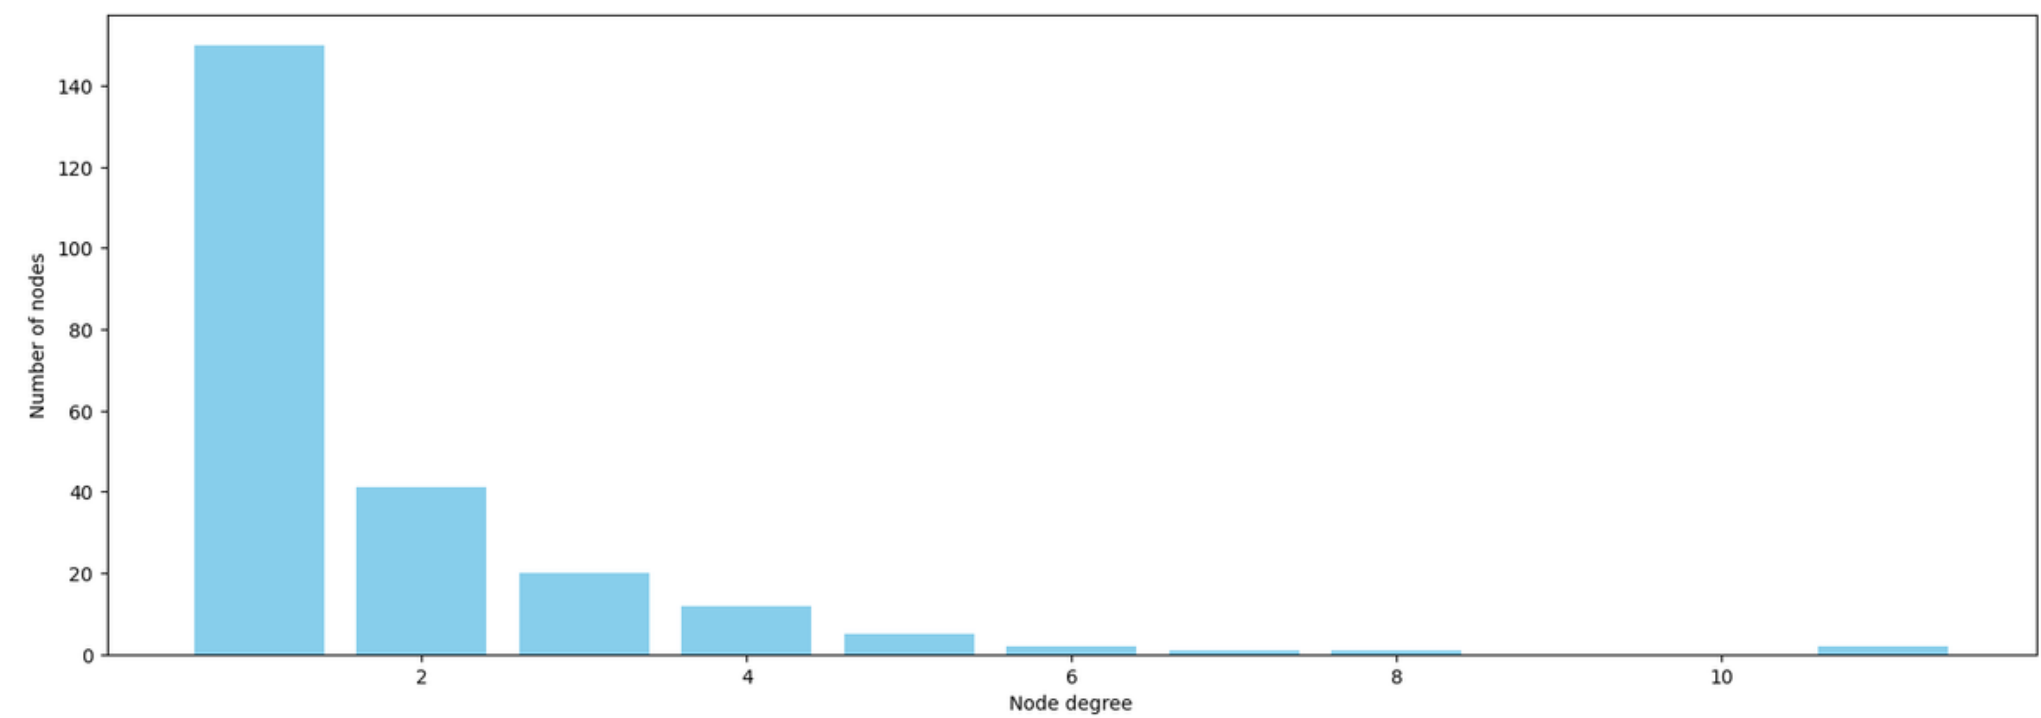
subgraph 0: Data(x=[500, 1433], edge_index=[2, 817], y=[500], train_mask=[500], val_mask=[500], test_mask=[500], n_id=[500], e_id=[817], input_id=[40], batch_size=40)
subgraph 1: Data(x=[563, 1433], edge_index=[2, 918], y=[563], train_mask=[563], val_mask=[563], test_mask=[563], n_id=[563], e_id=[918], input_id=[40], batch_size=40)
subgraph 2: Data(x=[688, 1433], edge_index=[2, 1324], y=[688], train_mask=[688], val_mask=[688], test_mask=[688], n_id=[688], e_id=[1324], input_id=[40], batch_size=40)
subgraph 3: Data(x=[230, 1433], edge_index=[2, 414], y=[230], train_mask=[230], val_mask=[230], test_mask=[230], n_id=[230], e_id=[414], input_id=[20], batch_size=20)

V Implementation

Degree



original graph degree



last subgraph degree

V Implementation

Accuracy calculation

```
1 def accuracy(pred_y, y):  
2 |     return((pred_y == y).sum()/len(y)).item()
```

GraphSAGE

```
1 class GraphSAGE(torch.nn.Module):  
2 |     def __init__(self, dim_in, dim_h, dim_out, agg = 'mean'):  
3 |         super().__init__()  
4 |         self.sage1 = SAGEConv(dim_in, dim_h, aggr=agg)  
5 |         self.sage2 = SAGEConv(dim_h, dim_out, aggr=agg)  
6 |         self.optimizer = torch.optim.Adam(self.parameters(),  
7 |                                           lr=0.01,  
8 |                                           weight_decay=5e-4)  
9 |  
10 |    def forward(self, x, edge_index):  
11 |        h = self.sage1(x, edge_index)  
12 |        h = torch.relu(h)  
13 |        h = F.dropout(h, p=0.5, training=self.training)  
14 |        h = self.sage2(h, edge_index)  
15 |        return F.log_softmax(h, dim= 1)  
16 |
```

$$\mathbf{h}'_i = \mathbf{W}_1 \mathbf{h}_i + \mathbf{W}_2 \cdot \text{mean}_{j \in \mathcal{N}_i}(\mathbf{h}_j)$$

V Implementation

GraphSAGE

training

```
17 def fit(self, data, epochs):
18     criterion = torch.nn.CrossEntropyLoss()
19     optimizer = self.optimizer
20
21     self.train()
22     for epoch in range(epochs+1):
23         total_loss = 0
24         acc = 0
25         val_loss = 0
26         val_acc = 0
27
28         # train on batches
29         for batch in train_loader:
30             optimizer.zero_grad()
31             out = self(batch.x, batch.edge_index)
32             loss = criterion(out[batch.train_mask], batch.y[batch.train_mask])
33             total_loss += loss.item()
34             acc += accuracy(out[batch.train_mask].argmax(dim=1), batch.y[batch.train_mask])
35             loss.backward()
36             optimizer.step()
37
38         # validation
39         val_loss += criterion(out[batch.val_mask], batch.y[batch.val_mask])
40         val_acc += accuracy(out[batch.val_mask].argmax(dim=1), batch.y[batch.val_mask])
41
42         # Print metrics every 10 epochs
43         if epoch % 5 == 0:
44             print(f'Epoch {epoch:>3} | Train Loss: {loss/len(train_loader):.3f} '
45                   f'| Train Acc: {acc/len(train_loader)*100:>6.2f}%'
46                   f'| Val Loss: {val_loss/len(train_loader):.2f} '
47                   f'| Val Acc: {val_acc/len(train_loader)*100:.2f}%')
48
```

test

```
49 @torch.no_grad()
50 def test(self, data):
51     test_acc = 0
52     self.eval()
53
54     for batch in train_loader:
55         out = self(batch.x, batch.edge_index)
56         test_acc += accuracy(out[batch.test_mask].argmax(dim=1), batch.y[batch.test_mask])
57
58     print(f'GraphSAGE test accuracy: {test_acc/len(train_loader)*100:.2f}%')
```

run code

```
1 exp1_max_small = GraphSAGE(dataset.num_features, 64, dataset.num_classes, agg='max')
2 print(exp1_max_small)
3
4 # Train
5 exp1_max_small.fit(data, 30)
6
7 # Test
8 exp1_max_small.test(data)
```


V Implementation

Result

Train loss


Mean

Mean-train loss



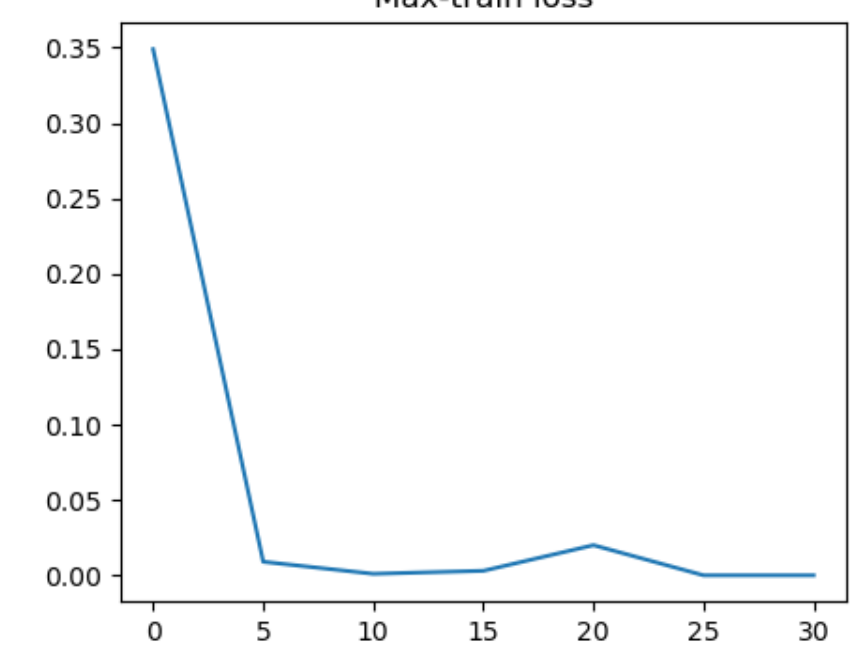
LSTM

LSTM-train loss




Max

Max-train loss




Validation loss


Mean-validation loss



LSTM-validation loss



Max-validation loss



V Implementation

Paper Result

Table 1: Prediction results for the three datasets (micro-averaged F1 scores). Results for unsupervised and fully supervised GraphSAGE are shown. Analogous trends hold for macro-averaged scores.

Name	Citation		Reddit		PPI	
	Unsup. F1	Sup. F1	Unsup. F1	Sup. F1	Unsup. F1	Sup. F1
Random	0.206	0.206	0.043	0.042	0.396	0.396
Raw features	0.575	0.575	0.585	0.585	0.422	0.422
DeepWalk	0.565	0.565	0.324	0.324	—	—
DeepWalk + features	0.701	0.701	0.691	0.691	—	—
GraphSAGE-GCN	0.742	0.772	0.908	0.930	0.465	0.500
GraphSAGE-mean	0.778	0.820	0.897	0.950	0.486	0.598
GraphSAGE-LSTM	0.788	0.832	0.907	0.954	0.482	0.612
GraphSAGE-pool	0.798	0.839	0.892	0.948	0.502	0.600
% gain over feat.	39%	46%	55%	63%	19%	45%

Implementaion Result

Mean	
small	big
65.85%	65.86%
LSTM	
small	big
59.27%	64.00%
Max	
small	big
63.44%	62.90%

2023/08/01

DSAIL summer-internship

STUDENT

유혜원

Hyewon Ryu

Thank you
for listening